

What Bayesforce Does: Strategy, Capabilities & Delivery Model

LAYER 1 STRATEGY (Worldview & Principles)

Why Bayesforce exists, the problem we own, and the unchanging principles that govern how we work.

I. Vision

Systems should carry more of the machinery of an organization so its people can carry more of its ambition.

Organizations exist to amplify what people can accomplish. Yet as they grow, an increasing share of human capacity is consumed by the machinery of running the organization: moving information, coordinating handoffs, reconciling data, and keeping processes moving.

Bayesforce envisions organizations where intelligent systems carry more of this machinery, while people spend more of their capacity on judgment, decisions, relationships, and growth.

The goal is not fewer people. **It is more capable organizations.**

II. Mission

To make organizations move as fast as the world around them.

AI is advancing faster than organizations can absorb it. But businesses cannot continually stop, replace their systems, retrain their teams, and rebuild their operations every time technology changes.

Bayesforce makes AI part of the existing operating fabric of an organization.

We connect the systems and information businesses already depend on, engineer AI into the workflows that run the business, and build the controls and capabilities required to operate them reliably.

Organizations should be able to adopt new capabilities without having to rebuild themselves.

III. Brand Promise

Make More Happen.

Every organization has work that consumes capacity without creating enough value.

People chase information. Systems remain disconnected. Teams reconcile data. Decisions wait for inputs. Work gets handed from one person to another simply because the machinery underneath it cannot carry it forward.

We call this operational drag.

Bayesforce exists to remove it.

We find where operational drag is holding an organization back and engineer AI into those workflows so work can move further with less human coordination.

The result is simple:

Less operational drag. More organizational capacity.

IV. Philosophy

Technology changes. Our principles do not.

1. Work Before Technology

We start with the work, not the technology.

We understand how work actually moves through an organization: what triggers it, what information it requires, where judgment is involved, where systems interact, and where it gets stuck.

The work determines the solution. We do not begin with a technology looking for a problem.

2. Make Technology Invisible

Technology should fit into the organization, not ask the organization to fit around it.

Bayesforce works through the systems, tools, and habits people already rely on. We add intelligence behind existing workflows wherever possible rather than introducing another layer of software for people to learn and manage.

The technology can be sophisticated underneath. **The experience should remain simple.**

3. Earn Autonomy

AI does not get autonomy because it is technically possible. It earns autonomy through evidence.

As a workflow becomes more reliable, observable, controlled, and well-understood, the system can take on more responsibility. We automate as much as is sensible for the work and its consequences, while keeping people involved wherever judgment or accountability requires it.

Autonomy is earned, not assumed.

4. Prove the Delta

Deployment is not success.

Every Bayesforce capability should have a meaningful baseline and a measurable outcome: faster processing, higher throughput, fewer errors, lower intervention, lower cost, or greater capacity.

We care about the difference the system creates in the real operation of the business.

If we cannot show the delta, we have not finished the job.

5. Design for Ownership

We build capability, not dependence.

Bayesforce designs systems that customers can understand, operate, govern, and improve. We work alongside internal teams, transfer knowledge, document important decisions, and build for long-term ownership.

The customer should become **more capable because Bayesforce was there.**

LAYER 2

BUSINESS DEFINITION (Capability & Portfolio)

Translating our worldview into concrete capabilities and the systems required to put them to work.

V. What We Do

Bayesforce is an AI capability engineering firm.

We identify high-friction organizational workflows, connect the systems and data involved, provide AI with the necessary enterprise context, and engineer intelligent systems directly into existing operations.

The workflow is our unit of engagement. We do not sell AI as a standalone layer; we engineer a production capability around a specific piece of organizational work.

We are fundamentally **function-agnostic but workflow-specific**.

The common denominator across our engagements is never the department—it is the nature of the work itself. Wherever repetitive knowledge work involves distributed information, context gathering, judgment, handoffs, and system interactions, Bayesforce engineers AI to carry the machinery forward.

The Unit of Work & Customer Output

A Bayesforce engagement produces a production-ready operational capability around a defined workflow: connected data and context, AI-driven execution, appropriate controls and human review, measurable performance, and the ability for the customer's team to operate and improve it.

VI. Capability Architecture

Bayesforce operates across four core capabilities. Together, these capabilities span what it takes to put AI to work inside an organization: establishing the foundations, executing the work, ensuring it can be trusted in production, and transferring the capability to the people who will own it.

CAPABILITY	CORE FOCUS & ROLE IN THE ENTERPRISE
1. AI Data Engineering	Foundational data pipelines, source system integration, and enterprise context engineering ("Second Brain").
2. AI Coworkers	Central execution engine building digital coworkers across the customer's existing tech stack.
3. AI Trust & Governance	Production-readiness layer ensuring AI is evaluated, observed, cost-optimized, and secured.
4. AI Enablement	Institutional capability transfer, technical training, and operational self-sufficiency.

1. AI Data Engineering

AI Data Engineering builds the foundational data plumbing and context mechanisms required for AI systems to operate accurately within an enterprise.

This capability comprises two core disciplines:

A. Data Pipelines (Integration & Preparation)

Connects source systems and prepares enterprise data for AI consumption:

- **Enterprise Source Systems Integration:** Direct connectors to ERPs, CRMs, ticketing systems, relational/NoSQL databases, cloud object storage, and internal communication tools.
- **End-to-End Pipeline Engineering:** Ingestion, normalization, schema transformation, and continuous real-time synchronization.
- **Data Quality & Reconciliation:** Schema validation, automated data consistency checks, missing-field resolution, and error handling.
- **AI-Ready Datasets:** Curating structured, clean, and pre-processed data streams optimized for agent consumption.

B. Context Engineering (Organizational Memory & Retrieval Architecture)

Constructs the contextual intelligence and retrieval infrastructure ("Second Brain") that models require to understand enterprise realities:

- **Document Understanding:** Parsing, extracting, and structuring messy, unstructured documents (PDFs, multi-format contracts, invoices, compliance records, technical manuals).
- **Contextual & Semantic Retrieval:** Hybrid retrieval systems combining vector semantic search, lexical matching, and metadata filtering.
- **Organizational Memory & Knowledge Graphs:** Mapping relationships between entities, accounts, transactions, and historical communications into persistent memory graphs.
- **Structured Context Feeding:** Dynamic assembly of minimal, high-signal context windows for downstream agents.

2. AI Coworkers (The Central Execution Engine)

AI Coworkers is the primary execution engine of Bayesforce. We design, build, and deploy autonomous digital coworkers that execute complex, multi-step operational work directly across the customer's existing tech stack.

Rather than introducing standalone chat interfaces or siloed software, digital coworkers operate behind the scenes and through the tools teams already use (Slack, Teams, CRMs, ERPs, email, databases).

Core Capabilities

- **Multi-System Tool Orchestration:** Connecting reasoning models to deterministic APIs, database queries, and enterprise applications to perform real operational actions.
- **Context-Aware Reasoning Engines:** Equipping digital coworkers with relevant schemas, playbooks, guidelines, and memory to execute domain-specific tasks.
- **Human-in-the-Loop Collaboration:** Designing intuitive approval checkpoints, escalations, and review boundaries for high-consequence edge cases.
- **Resilient Execution & Error Handling:** Hardening agent runtimes with retry logic, fallback behaviors, and graceful degradation when upstream tools fail.

3. AI Trust & Governance

AI must not only work in prototypes and demos; it must be trustworthy enough to operate in production. AI Trust & Governance provides the enterprise control layer across four distinct pillars:

A. AI Evaluations (Evals & Benchmarks)

- **Workflow-Specific Evals:** Custom evaluation datasets reflecting the company's real edge cases and domain criteria.
- **Agent Performance & Correctness Metrics:** Quantitative measurement of task completion, tool-call accuracy, and reasoning fidelity.
- **Regression Test Suites:** Automated CI/CD testing to prevent model performance degradation when prompts, data, or base models change.
- **Outcome-Based Quality Benchmarks:** Benchmarking AI output quality directly against human expert standards.

B. AI Observability (Tracing & Monitoring)

- **End-to-End Tracing:** Granular visibility into agent execution paths, retrieved context, decision inputs, tool calls, outputs, and execution traces.
- **Failure & Latency Tracking:** Real-time monitoring of API timeouts, token bottlenecks, tool failures, and fallback triggers.
- **Root-Cause Analysis:** Fast debugging mechanisms to pinpoint why a workflow failed or hallucinated.
- **Production Behavior Dashboards:** Live telemetry on workflow throughput, error rates, and human intervention frequency.

C. AI FinOps (Cost Optimization & Unit Economics)

- **Workflow-Level Cost Analysis:** Granular tracking of compute, API, and token expenditure per workflow and business unit.
- **Cost Optimization Architectures:** Implementing intelligent caching, prompt compression, and dynamic model routing.
- **AI Bill Management & Predictability:** Forecasting operational AI costs and eliminating runaway compute spend.

D. AI Safety & Governance (Compliance & Access Control)

- **Granular Permissions & Role-Based Access:** Ensuring agents only access data and tools permitted for specific operational roles.
- **Data Protection & PII Leakage Controls:** Automated redaction of sensitive personal data, financial details, and proprietary intellectual property.
- **Prompt Injection Defense & Red-Teaming:** Hardening agentic endpoints against adversarial attacks, jailbreaks, and unauthorized tool calls.

- **Policy Enforcement & Audit Logging:** Auditable logs of agent actions, tool calls, system changes, and human approvals for regulatory and compliance audits.

4. AI Enablement (Capability Transfer)


For organizations that want to build, operate, and maintain AI systems internally, AI Enablement provides structured capability and knowledge transfer.

Rather than offering generic training or passive corporate L&D, this capability ties enablement directly to tangible operational impact.

Core Offerings

- **Enterprise AI Literacy:** Baseline education across leadership and functional teams on applied AI capabilities, limitations, and operational opportunities.
- **AI Enablement for Leaders & Managers:** Executive frameworks for identifying high-ROI automation wedges, managing human-agent teams, and designing governance policies.
- **Workflow-Specific Technical Training:** Hands-on prompt engineering, agent orchestration, and eval design training for internal technical and operations teams.

The 4-Step Capability Transfer Model



VII. Workflow Scope & Portfolio

Bayesforce is not organized around a fixed set of departments or industries. We operate wherever organizations have recurring work that can be carried forward by better use of data, context, reasoning, and systems.

Our commercial scope is therefore defined at two levels:

- **Business Function:** a recurring area of organizational work that may span multiple departments or teams.
- **Operational Workflow:** a specific, repeatable sequence of activities within that function that produces a defined outcome.

For example, Revenue Operations is a business function; lead qualification, CRM enrichment, pipeline inspection, and sales follow-up are operational workflows within it.

This distinction matters because organizations are structured around people and reporting lines, while work often crosses those boundaries.

Bayesforce is fundamentally function-agnostic but workflow-specific.

We do not automate a department for the sake of automating the department. We identify the workflows within it where operational drag is significant enough, the underlying systems are accessible enough, and the resulting improvement is valuable enough for AI to carry a meaningful share of the work.

Supported broadly. Pursued selectively.

Bayesforce can operate across a broad range of business functions, but broad capability does not imply broad pursuit. Our portfolio remains deliberately selective, based on workflow economics, technical feasibility, risk, repeatability, and strategic fit.

1. Where Bayesforce Operates

Bayesforce supports operational workflows across eight broad business functions.

BUSINESS FUNCTION	REPRESENTATIVE WORKFLOW FAMILIES	STRATEGIC POSITION
Revenue Operations	Lead qualification, CRM enrichment, pipeline inspection, sales follow-up, commercial research, marketing operations	CORE
Customer Operations & Service Delivery	Ticket investigation, resolution, escalation, customer history, onboarding, cross-team service coordination	CORE
Finance & Accounting	Invoice processing, reconciliation, exception handling, cash application, collections, reporting	CORE
Procurement & Vendor Operations	Purchase requests, policy verification, vendor onboarding, quote comparison, contract-to-PO validation	CORE
Business Operations & Shared Services	Internal requests, email-to-workflow, SOP execution, approvals, exception management, cross-team coordination	CORE
Legal, Risk & Compliance	Contract intake, obligation tracking, compliance questionnaires, audit evidence, policy checking	EXPANSION
Technology Operations	IT service desk, access provisioning, incident triage, operational knowledge resolution	EXPANSION
People Operations	Employee onboarding, HR service desk, policy execution, credential verification, employee lifecycle workflows	EXPANSION

These are not separate Bayesforce products. They are **places where the same underlying workflow capabilities can be applied to different forms of organizational work**.

Marketing is included selectively within Revenue Operations where the work is operational rather than primarily creative: campaign operations, lead handoff, attribution reconciliation, reporting, and commercial research are relevant; generic content generation is not a strategic Bayesforce focus.

Likewise, Bayesforce may operate around technology teams, but does not position itself as a developer-tooling or software-development infrastructure company. The focus is the **operational work surrounding technology**, such as service operations, access management, incident coordination, and evidence gathering.

2. The Workflows We Are Built to Support

Although the business functions differ, the workflows that fit Bayesforce tend to share the same underlying structure.

Our strongest targets typically contain several of the following characteristics:

- **High operational drag:** substantial human time is spent moving information, coordinating steps, checking work, or chasing exceptions.
- **High recurrence:** the workflow happens frequently enough for improvement to compound.
- **Multiple systems:** the work crosses CRMs, ERPs, ticketing systems, databases, documents, email, spreadsheets, or internal tools.
- **Messy information:** important inputs are distributed across unstructured or semi-structured sources.
- **Mixed reasoning:** the workflow combines deterministic business rules with contextual interpretation or judgment.
- **Meaningful handoffs:** progress depends on several systems, teams, or people coordinating correctly.
- **Measurable outcomes:** the workflow has a clear baseline and a meaningful delta can be measured.
- **Recoverable consequences:** automation can operate within appropriate approval, exception, and rollback boundaries.
- **Reusable architecture:** the underlying pattern can be applied again rather than becoming a one-off bespoke system.

This leads to a simple strategic observation:

The best Bayesforce workflow is not the one with the most AI. It is the one where AI can remove meaningful operational drag with acceptable risk and measurable impact.

3. Our Feasibility Screen

Not every workflow that can be automated should be automated.

Before committing Bayesforce resources to a workflow, we evaluate it across eight dimensions:

DIMENSION	HIGH-FIT CONDITION	LOW-FIT CONDITION
1. Operational Drag	Significant recurring manual effort, coordination, delay, or exception handling	Minor inconvenience or negligible manual burden
2. Frequency	Daily or continuously recurring work with meaningful volume	Rare, episodic, or one-off work
3. Economic Impact	Clear effect on revenue, cash, capacity, turnaround time, cost, or risk	Soft or difficult-to-measure benefit
4. AI Suitability	Requires contextual understanding, extraction, synthesis, classification, reasoning, adaptive routing, or interaction with ambiguous information	AI adds little material value because the workflow is already simple, deterministic, and efficiently handled by conventional software or automation
5. System Connectivity	Relevant systems expose APIs, databases, events, documents, or reliable interfaces	Poorly accessible or fundamentally disconnected sources
6. Consequence & Control	Errors are recoverable and human approval can sit at consequential boundaries	Irreversible or unacceptable consequences without adequate controls
7. Repeatability	Common underlying pattern that can be reused across workflows or customers	Highly bespoke process with little architectural reuse
8. Commercial Potential	Sufficient recurring value to support meaningful Bayesforce ownership	Commodity, low-ticket, or one-off consulting economics

These dimensions should not be treated as a mechanical score that decides whether a workflow lives or dies.

They are a **strategic filter**.

A workflow with high value and high risk may still be attractive, but it requires stronger controls, narrower autonomy, deeper domain expertise, and a higher trust threshold. A workflow with low risk but little economic value may be technically easy but commercially irrelevant.

The question is therefore not simply:

Can AI do this?

It is:

Should Bayesforce make AI responsible for this work, and does doing so create enough value to matter?

4. Strategic Portfolio Prioritization

The feasibility screen produces three broad categories of work.

TIER 1 Core Target Portfolio

These workflows should form the center of Bayesforce's initial commercial portfolio because they align most directly with the operating thesis: high recurring friction, multiple systems, strong measurability, meaningful human coordination, and clear opportunities for AI-assisted or autonomous execution.

Customer Operations & Service Delivery

Representative targets:

- Context-driven ticket investigation
- Tier-1 resolution
- Customer history reconstruction
- Cross-team service checks
- Escalation packet preparation
- Customer onboarding coordination

Why it fits: Customer operations combines high transaction volume, rich system connectivity, significant context requirements, and clear operational metrics such as resolution time, backlog, escalation rate, and human intervention.

Finance & Accounting

Representative targets:

- Invoice intake and extraction
- Purchase order and invoice reconciliation
- Invoice exception handling
- Cash application
- Vendor statement reconciliation
- AR dispute management
- Collections prioritization

Why it fits: Finance workflows often contain substantial repetitive coordination, multiple source systems, structured approval boundaries, and measurable economic outcomes. They also reward strong reconciliation, evidence, auditability, and human approval architecture.

Business Operations & Shared Services

Representative targets:

- Internal request desks
- Email-to-workflow automation
- SOP execution
- Approval routing
- Cross-team handoff coordination
- Exception management
- Recurring reporting and decision packs

Why it fits: This is the purest expression of the Bayesforce thesis. Much of organizational drag exists in the connective tissue between systems and teams rather than inside any single application.

Procurement & Vendor Operations

Representative targets:

- Purchase request intake and policy verification
- Vendor onboarding
- Credential and document verification
- Quote comparison
- Contract-to-PO validation
- Vendor compliance workflows

Why it fits: Procurement combines documents, rules, systems, approvals, external parties, and repeated reconciliation. The work is operationally important while retaining clear control boundaries.

Revenue Operations

Representative targets:

- Lead qualification and routing
- Account and CRM enrichment
- CRM hygiene
- Pipeline inspection
- Sales follow-up orchestration
- Account research and briefing

Why it fits: Revenue workflows are data-rich, highly connected, frequent, and relatively easy to measure. They provide clear commercial outcomes, although Bayesforce should remain focused on operational execution rather than generic "AI sales" or content-generation propositions.

TIER 2 Strategic Expansion Portfolio

These workflows are attractive but should generally follow successful deployments in the core portfolio. Their constraints may include higher consequences, greater domain specialization, longer trust cycles, or heavier dependence on workflow-specific controls.

Legal, Risk & Compliance

Potential expansion areas:

- Contract intake and playbook triage
- Obligation tracking
- Vendor compliance
- Compliance questionnaire response
- Audit evidence collection
- Policy compliance checking

The strongest opportunities are those where AI can gather, extract, compare, monitor, and prepare while authorized professionals retain consequential approval.

Technology Operations

Potential expansion areas:

- IT service desk
- Access provisioning
- Incident intake and context triage
- Operational knowledge resolution
- Release-readiness and evidence workflows

The strategic boundary is important:

Bayesforce automates operational work around technology; it does not need to become the technology stack for software developers.

People Operations

Potential expansion areas:

- Employee onboarding and offboarding
- Internal HR service operations
- Policy question resolution and action execution
- Credential and document verification
- Employee lifecycle administration

Administrative and service workflows are generally more compatible with the Bayesforce model than autonomous decisions about hiring, promotion, compensation, or other high-consequence employment outcomes.

TIER 3 Deprioritized or Non-Core

Some workflows are technically feasible but strategically weak for Bayesforce.

We should be deliberate about not building a business around them.

Generic Content Generation

Copywriting, generic social posts, generic marketing assets, and similar work are increasingly commoditized and do not inherently involve the systems, coordination, and operational machinery at the center of the Bayesforce thesis.

Standalone Enterprise Knowledge Retrieval

"Chat with your documents" can be useful, but a knowledge interface by itself is not enough. Bayesforce should preferentially connect knowledge to a workflow that leads to a decision, action, system update, or measurable operational outcome.

Generic AI Training

AI literacy and broad awareness programmes can be valuable, but they should generally be attached to a real implementation or capability-transfer engagement rather than become a standalone training business.

Generic Project and Status Summaries

Summarization alone does not create a sufficiently strong Bayesforce wedge when the underlying workflow, decision, or execution remains unchanged.

Developer Infrastructure

Bayesforce should avoid competing directly as a general-purpose developer platform, coding agent, CI/CD product, or software engineering infrastructure provider.

Highly Autonomous High-Consequence Decisions

Workflows involving irreversible decisions, material legal exposure, significant financial exposure, or other severe consequences should not become early Bayesforce targets unless the workflow can be decomposed into well-controlled stages with explicit human accountability.

5. The Enabling Layer Is Cross-Functional

Data, context, trust, governance, and capability transfer do not belong to one business function.

They are the **shared infrastructure that allows Bayesforce to operate across all of them.**

The following capabilities therefore sit across the entire workflow universe:

- Source-system integration
- Data pipelines and synchronization
- Document understanding and normalization
- Context assembly and retrieval
- Evaluations and regression testing
- Observability and exception monitoring
- Security, access control, and auditability
- Cost and workflow economics
- Capability transfer and internal enablement

These capabilities may be sold independently in selected circumstances, but strategically they are strongest when attached to a real operational workflow.

Data, context, and control are not separate from the workflow. They are what make it capable of operating reliably in production.

6. The Strategic Wedge: How We Enter a Customer

Bayesforce should resist entering a market by saying:

"We automate Finance."

or:

"We build AI for Customer Support."

Those statements are too broad.

The initial commercial wedge should instead look like:

We remove a specific piece of operational drag from a specific workflow.

Our preferred commercial wedge is a single, high-leverage workflow—such as invoice exception handling or ticket investigation—expanding outward as Bayesforce proves value, gains system access, and establishes operational trust.

This creates a deliberate land-and-expand trajectory:

Workflow → Adjacent Workflows → Business Function → Cross-Functional Operations

The objective is not to sell a collection of unrelated automations.


It is to become increasingly embedded in the operational machinery of the organization.

7. Reusable Workflow Archetypes

Although organizations contain hundreds of distinct workflows, many are variations of a much smaller number of underlying operational patterns.

This is strategically important.

Bayesforce should not rebuild every workflow from zero. We should develop reusable architecture around recurring patterns and then contextualize those patterns to the customer's systems, data, policies, and operating environment.



8. The Portfolio Principle: How We Select What to Pursue

The workflow universe is intentionally broad. The Bayesforce commercial portfolio is deliberately narrow.

While our **Strategic Wedge** defines how we expand within a customer, the **Portfolio Principle** governs how we allocate company resources across the market:

Remain broad in where we can operate, but narrow in what we choose to pursue.

We prioritize workflows where operational drag is large, recurrence is high, systems are accessible, outcomes are measurable, risks are controllable, and the underlying architecture can compound across customers.

We do not chase every workflow that AI can touch.

We build depth where the same underlying machinery can repeatedly remove meaningful operational drag.

Many functions. Many workflows. A small number of repeatable patterns.

That is the basis on which Bayesforce can remain function-agnostic without becoming strategically unfocused.

LAYER 3 DELIVERY (The Operational Methodology)

The unified 11-stage delivery model through which every workflow is engineered, governed, calibrated, and transferred.

VIII. The Bayesforce Workflow

Every workflow we engineer is mapped to the same underlying lifecycle, with stages applied according to the nature of the work and the level of autonomy it warrants.

This 11-stage delivery model is not merely an AI agent architecture; it is Bayesforce's end-to-end delivery methodology. It brings together data plumbing, contextual grounding, autonomous reasoning, deterministic system actions, human governance, and empirical measurement into a unified operational discipline.

The 11-Stage Workflow Lifecycle



- 1 **Trigger / Intake:** An operational event initiates the workflow (e.g., incoming email, webhook, file upload, scheduled job, status update).
- 2 **Ingest & Collect:** Gathers the relevant operational payload, source files, attachments, and transaction metadata.
- 3 **Understand & Classify:** Parses incoming intent, extracts structured entities, categorizes the request, and determines operational routing.
- 4 **Context Retrieval:** Pulls relevant institutional memory, customer history, business playbooks, and schema definitions from the context layer.
- 5 **Validate & Verify:** Checks data completeness, constraints, business prerequisites, and schema integrity before acting.
- 6 **Reason & Decide:** Performs probabilistic reasoning, trade-off analysis, and decision formulation based on enterprise guidelines.
- 7 **Orchestrate & Execute:** Coordinates tool calls and executes deterministic actions across target software tools (CRMs, ERPs, APIs).
- 8 **Output Generation:** Drafts the required operational output (e.g., email reply, contract redline, payment batch, reconciliation report).
- 9 **Human Review & Approval:** Human-in-the-loop checkpoint where an authorized operator reviews, modifies, or signs off on high-stakes outputs.
- 10 **System of Record Update:** Commits the finalized state, transaction logs, and updates to primary databases and ERP/CRM systems.
- 11 **Outcome Measurement & Learning:** Compares actual outcome against the expected baseline, measures delta, and feeds learnings back into prompt, retrieval, and model tuning.

Mapping Principles to the Workflow

Our core philosophy directly governs each phase of the workflow. These mappings are illustrative rather than exclusive; each principle governs the workflow end-to-end, but becomes especially visible in the stages below:

PRINCIPLE	ROLE IN THE WORKFLOW	PRIMARY STAGES
1. Work Before Technology	Understand the operational triggers, information requirements, and actual friction before applying AI.	Stages 1–3 (Trigger, Ingest, Understand)
2. Make Technology Invisible	Work through existing systems, tools, and habits rather than adding separate software layers.	Stages 4, 7, 10 (Context, Execute, Record)
3. Earn Autonomy	Validate prerequisites, establish human review boundaries, and scale agent autonomy through evidence of reliability.	Stages 5, 6, 9 (Validate, Reason, Review)
4. Prove the Delta	Continuously measure real-world performance, speed, accuracy, and business impact against pre-AI baselines.	Stage 11 (Outcome Measurement & Learning)
5. Design for Ownership	Equip internal teams to understand, operate, govern, and improve the workflow long-term.	End-to-End Governance & Capability Transfer

Sensible Autonomy Calibration

Not every workflow executes all 11 stages autonomously. Some steps are purely deterministic business logic; other steps require probabilistic LLM reasoning; critical edge cases remain entirely human-decided.

The objective is **not maximizing autonomy for its own sake**, but automating as much of the workflow as is sensible, reliable, and safe for the organization and its risk profile.

BAYESFORCE

AI Capability Engineering • Work Before Technology • Earn Autonomy • Prove the Delta

END OF STRATEGY, CAPABILITIES & DELIVERY MODEL SPECIFICATION